

Benchmark Gains Are Not Governance Gains

Paired, Budget-Matched Evaluation of Small Safety Guards Under Domain Shift

ANONYMOUS AUTHOR(S)

A *safety guard* — a small classifier that labels an incoming request safe or unsafe before a language-model system acts on it — is the mechanism by which a written content policy becomes an enforced decision. Guards are chosen by a rule that is itself never audited: fine-tune a small open model on available moderation data, score it on a public suite, ship the best checkpoint. We treat that rule as the object of study and ask what a benchmark gain is valid *for*.

We evaluate four small open checkpoints (4×5 seeds) under a protocol with four commitments: each tuned guard is compared only with its *own* pre-tuning checkpoint on identical rows; sources named in the training manifest are separated from sources held out at the source level; recalls are read only at *matched false-alarm budgets*; and three tiers of evidence are never pooled. The headline changes under each. Fine-tuning raises macro average precision on represented sources by +0.3234 [+0.2647, +0.3690] while held-out transfer moves -0.0589 [-0.0837, -0.0321], pooling per-checkpoint effects of opposite sign. Read at each base guard’s own alarm rate, the apparent transfer-recall gain *reverses* on all four checkpoints ($0.517 \rightarrow 0.185$), and recall on hard attacks falls $0.780 \rightarrow 0.171$. Restricted to the alarm budget an inline guard can afford, the same comparison is $3.0\times$ larger; and the sourcing conclusion — self-host or buy a hosted frontier model — *inverts* between traffic a manifest represents and traffic it does not.

A frozen, dual-labeled mortgage benchmark then isolates the stratum a general-safety score cannot see: requests that read as ordinary professional prose while soliciting redlining by proxy or a pretextual adverse-action reason. All four guards rank one such request *below* the median benign inquiry in the same split, and the protected-attribute invariance gate we built to check them fails its own robustness test — reported here as a negative result about instrument design. Our contribution is an evidence discipline, not a new guard: four disclosures, computable from per-row scores an operator already holds, that make specialization, threshold drift, and domain blindness legible before a benchmark number is allowed to license a deployment claim.

Additional Key Words and Phrases: safety guards, content moderation, evaluation validity, distribution shift, algorithmic auditing, fair lending, measurement

1 INTRODUCTION

Between a user’s request and a language-model system’s response there is increasingly a second, much smaller model whose only job is to decide whether the request may proceed. These *guards* — Llama Guard, WildGuard, ShieldGemma, Granite Guardian, Qwen3Guard, and a long tail of in-house fine-tunes [26, 31, 46, 47, 59] — are the point at which a written policy stops being a document and becomes an enforced decision about a particular person’s particular request. They are governance infrastructure in the sense Gorwa et al. [24] and Gillespie [23] describe for platform moderation: small, cheap, invisible, and consequential.

How does an operator choose one? Overwhelmingly by a rule that is never itself audited: take a compact instruction model, attach a low-rank adapter, fine-tune on whatever moderation data is at hand [18, 29], score the result on a public guardrail suite [6], and keep the best-scoring checkpoint. The rule is cheap, reproducible, and — as an argument that the resulting system is fit to screen real traffic — close to vacuous.

Our claim is not that fine-tuning harms guards. It is narrower: **a benchmark gain measures something, and that something is usually not what the deployment decision needs to know.** Following the measurement-modelling tradition in this venue [8, 32, 48], we ask what construct a guard score operationalizes, and find four places where the operationalization and the deployment claim come apart. Each is demonstrated on a fixed panel of four small open checkpoints, and each is a *reporting* failure before it is a modelling one:

- (1) **Whose sources?** The score rises sharply on the corpora the guard’s own training manifest names and does not, on average, move on corpora held out at the source level: it measures acquaintance with the instrument, not the construct (Section 4).
- (2) **At whose alarm rate?** The apparent off-source recall gain is bought by alarming far more often; equalize the false-alarm budget and it reverses on every checkpoint, hardest on the hardest attacks, and restricting the metric to the region an inline guard operates in roughly triples the cost (Section 5).
- (3) **On whose traffic?** Against a hosted frontier model on identical rows at a common budget, the small guards *lose* on unfamiliar prompts and *win* on sources their manifest names — the ordering is a property of the traffic regime, not of the model, so one number cannot support a sourcing decision (Section 6).
- (4) **Against whose policy?** In a regulated domain the dangerous request is often fluent, polite, procedurally framed text with no jailbreak, slur, or injection, for which a general-safety taxonomy has no label; we isolate that stratum with a dual-labeled, HMDA-grounded mortgage benchmark and show a worked miss (Section 7).

Figure 1 is those four findings in one spread. The distributive stakes deserve stating at the outset, because the aggregates hide them: a guard’s false positives fall on legitimate users whose requests are blocked, its false negatives on whoever is harmed when a request is honored. On this panel fine-tuning moves *both* the wrong way off-source at once — pooled false alarms on held-out negatives climb 4.3% → 17.0% while recall on HarmBench attacks falls 78.0% → 60.0%. That is not a trade-off between two populations but a loss for both, invisible to the represented-source number that would have justified the deployment.

Contributions. (i) A **validity-first evaluation protocol** built from four commitments — strict pairing, source-level regime separation, matched alarm budgets, non-pooled evidence tiers — stated so a reader can check whether any published guard result satisfies them (Section 3). (ii) The **demonstration that each commitment is load-bearing**: dropping any one changes the sign or magnitude of the headline. (iii) A **dual-labeled regulated-domain instrument**, a worked end-to-end miss, and an honest negative result about the protected-attribute gate we built alongside it (Section 7). (iv) A **disclosure schedule** mapping each finding to the artifact an auditor, procurement reviewer, or regulator needs to read a guard benchmark correctly (Section 8).

What we do not claim. We introduce no model, metric, or training algorithm. The main analysis is retrospective and estimation-only on a purposively chosen panel inspected during development; intervals are conditional on that panel and are not hypothesis tests. The mortgage labels are model-assigned against written policy cards, *not* adjudicated by counsel or subject-matter experts. Nothing here is a causal claim, a universal claim about fine-tuning, a deployment recommendation, or a fair-lending finding (Section 9, Table 4).

2 GUARDS AS GOVERNANCE INFRASTRUCTURE

2.1 The object, and the inference under audit

The guards we study are prompt classifiers. Given a request x , a compact decoder-only model is prompted with a fixed template and its verdict read as a single-token margin — the difference of two decision logits, one forward pass, no generation:

$$s(x) = z_{\text{unsafe}}(x) - z_{\text{safe}}(x). \quad (1)$$

A threshold τ turns $s(x)$ into block or allow, and everything downstream of that comparison — who is refused service, which violating request proceeds — is decided by the pair (s, τ) . The evaluation literature reports mostly s .



Fig. 1. **Four ways a guard's benchmark score fails to answer the deployment question.** (a) Paired change in macro average precision against each guard's *own* pre-tuning checkpoint on identical rows: large and uniform on the corpora the training manifest names, near zero and split in sign on corpora held out at the source level. (b) Held-out recall at three thresholds; the middle bar is the tuned guard at its own calibrated cutoff — the comparison a leaderboard reports — where it alarms roughly four times as often on held-out negatives as its base. (c) The same comparison against a hosted frontier reference at a matched 5% budget points in opposite directions on the two traffic regimes (upper bar: a post-hoc summary over 3 purposively chosen corpora; lower: a paired comparison on external expert-annotated rows). (d) The same checkpoints, zero-shot, ranking a mortgage-policy violation label against the chance floor its own prevalence sets. Every plotted value is parsed by `figures/make_facct_figures.py` from the committed artifacts the tables `\input`, so no panel can drift from its table.

Two features of the setting are routinely elided. The errors are *asymmetric in incidence*, not merely in cost — a false positive is borne by a user with a legitimate request, a false negative by whoever the honored request harms — so one accuracy number aggregates across populations with no common welfare denominator, the objection Selbst et al. [55] raise to abstraction boundaries drawn for modelling convenience. And a guard is a *screening layer in front of a decision*, exactly where Raji et al. [49] argue functionality rather than fairness is the first thing to interrogate: a guard that does not work is not a fairness problem to mitigate but a non-functional component whose deployment was licensed by a bad measurement.

The selection rule's implicit inference — that a higher score is evidence of a better screening layer — must cross four gaps: from the benchmark's sources to real traffic; from a threshold-free ranking metric to a specific operating point; from a balanced pool to a real prevalence of unsafe requests; and from a general-safety taxonomy to the policy the operator must enforce. Each is known in isolation; what is missing is a protocol that measures all four *on the same rows for the same checkpoint*, so the size of each gap can be set against the size of the reported gain. In Jacobs and Wallach [32]'s terms these are four threats to the validity of one inference — source, operational, regime, and construct validity — and Sections 4 to 7 take them in turn.

2.2 Related work

Three strands meet here, and [Appendix A](#) discusses each at length. *Evaluation validity*: benchmarks under-determine the claims made from them — general-capability framings outrun what a fixed dataset supports [48], fairness benchmarks carry construct failures [8], behavioural testing beats aggregate scores [10, 53], progress on a fixed test set need not survive a change of collection process [52, 57], and generative-AI evaluations sit at a distance from the outcomes they are invoked to predict [38, 56, 58]. Where model cards and datasheets [22, 30, 44] established that documentation is part of the artifact, audit frameworks [43, 50, 51] established who consumes it; [Section 8](#) is written in that idiom, and Kapoor and Narayanan [34] is our closest methodological precedent for treating “was the evaluation *source* represented in training?” as a first-class reported quantity. *Guard evaluation*: released guards [26, 31, 46, 47, 59] are typically reported against in- and out-of-distribution blocks, adapting a general model into a guard is standard [18], and adjacent work finds safety degradation tracking dataset similarity [28], policy overfitting [41], surface heuristics [37], evadability [7, 25], poor calibration [40], and domain fine-tuning of released guards [1, 27, 36]. *Policy-grounded and regulated-domain evaluation*: human-labelled policy-compliance benchmarks [13, 14], regulation-grounded data [60], the canonical account of facially neutral proxies [4], its measurement in consumer lending [5], and language models audited as underwriters [9]; our protected-pair design is the standard counterfactual token-fairness construction [21, 35], repurposed from decision correctness to score invariance. [Table 5](#) places the closest work against the axes we measure: no row does all of it, and almost every column has a row that does that column better than we do.

3 THE PROTOCOL AND ITS FOUR COMMITMENTS

C1: compare a guard only with itself. Leaderboard comparisons put model X ’s guard beside model Y ’s, confounding “the fine-tune helped” with “ X was a better starting point.” We hold the checkpoint fixed and compare every tuned guard with its *own* pre-tuning weights, on identical rows, through the identical head of [Equation \(1\)](#) and the identical tie-aware metric, reporting a within-checkpoint change $\Delta = \text{SFT} - \text{base}$. This is the only construction that attributes a movement to the intervention rather than to the luck of the starting point, and it is what makes the sign splits in [Figure 1\(a\)](#) interpretable rather than noise.

C2: separate represented from source-held-out.

Every guard and its base are scored in two disjoint regimes. *Represented* rows are held back from the three corpora the manifest names (toxicchat [39], prompt_injections, jailbreak_classification); *transfer* rows come from four corpora whose *sources* were never trained on [11, 26, 33, 54]; and two single-class stress sets isolate the failure directions, OR-Bench [17] for over-refusal and HarmBench [42] for missed catches. The row-level/source-level distinction is where guard papers most often equivocate: held-out *rows* from a represented corpus measure whether the guard learned that corpus’s label conventions, held-out *sources* whether it learned anything portable, and reporting the first as generalization is the leakage-shaped error Kapoor and Narayanan [34] catalogue, committed at corpus rather than row level. We audit it — a committed overlap audit finds no exact, normalized, n -gram-containment (≥ 0.80), or character-shingle (Jaccard ≥ 0.70) overlap between manifest and transfer suite, and no evaluation-family collisions — while noting that the *represented* splits are not clean in the same sense (1.6%–5.0% of each sits within Jaccard 0.70 of a training row), which qualifies the represented margins, not the transfer ones.

C3: read recalls only at matched alarm budgets. A recall compared at unequal false-alarm rates is not a comparison of discriminative power but of two different policies. We therefore report thresholded quantities twice: at each guard’s

Table 1. **Paired base→tuned macro average precision by regime**, reproduced verbatim from the committed generated artifact. *Rep* = held-back rows from the three corpora named in the manifest; *Tr* = four corpora held out at the source level. Intervals are two-sided 95% hierarchical bootstrap, conditional on this fixed panel.

Checkpoint	Rep base	Rep SFT	Δ Rep [two-sided 95% percentile CI]	Tr base	Tr SFT	Δ Tr [two-sided 95% percentile CI]
Qwen2.5-1.5B	0.6334	0.9878	0.3544 [0.2731, 0.4150]	0.8187	0.7798	-0.0389 [-0.0829, 0.0062]
SmolLM2-1.7B	0.4524	0.9806	0.5282 [0.4555, 0.5748]	0.7904	0.8304	0.0400 [0.0003, 0.0776]
SmolLM3-3B	0.6621	0.9751	0.3130 [0.2419, 0.3701]	0.9102	0.8234	-0.0869 [-0.1114, -0.0613]
Qwen3-4B	0.8855	0.9837	0.0981 [0.0545, 0.1479]	0.9438	0.7939	-0.1499 [-0.1963, -0.1050]
Fixed-panel aggregate	–	–	0.3234 [0.2647, 0.3690]	–	–	-0.0589 [-0.0837, -0.0321]

own calibration-selected cutoff, and with the tuned guard re-thresholded so its pooled false-alarm rate on the evaluated negatives is *at or below its own base’s*. Only the second licenses a sentence about catching more or less. One limit travels with every such number: the budget is located by a quantile of the *evaluated* negatives and recall measured on those same rows, so it reads two ROC curves at a common point but is *not* a threshold a production system could place, since placing it requires the labels the system is predicting (Appendix B.1).

C4: never pool evidence tiers. The paper carries three kinds of evidence and never averages them: *retrospective panel* evidence (estimation-only on a fixed panel inspected during development); *external expert-annotated* evidence from a third-party corpus [12]; and *LLM-judge, policy-card-consistent* evidence from the mortgage instrument, where a model assigns labels against written policy cards with self-consistency checks but no human adjudication and no agreement statistic. A further qualifier, *analysis-preregistered*, applies to one arm whose estimands and criteria were fixed in a committed registry before any score existed; it is not data-blind, so it qualifies a retrospective row rather than forming a tier. Table 4 is the ledger, and its rule is that no sentence here is stronger than the tier of its weakest input.

Panel, manifest, estimand. Four open instruction checkpoints spanning 1.5–4B parameters (Qwen2.5-1.5B, SmolLM2-1.7B, SmolLM3-3B, Qwen3-4B) are fine-tuned from an identical frozen manifest of 1,200 rows — 400 per represented corpus, balanced 200 safe / 200 unsafe within each, decontaminated, data order fixed by seed — with five training seeds per checkpoint, giving 20 guards. One low-rank adapter configuration is held constant across the panel [29]; only the adapter trains, so “base” means the same weights *before* the guard adapter, read through the identical head. Each regime is summarized by tie-aware macro average precision (equal weight per benchmark, so a large corpus cannot dominate), and each paired change carries a two-sided 95% interval from a hierarchical bootstrap that keeps every adapter tied to its own base, resamples near-duplicate evaluation families rather than bare rows, and respects the checkpoint→seed grouping. The estimand is deliberately narrow — the change *conditional on this fixed panel and this manifest* — and these are estimation intervals, not significance tests (Appendix B).

4 Q1: THE GAIN SITS ON THE SOURCES THE MANIFEST NAMES

On represented sources the gain is large and, more tellingly, *uniform in its destination*: whatever the base, every tuned guard lands near macro average precision 0.98 (Table 1). SmolLM2-1.7B climbs from 0.4524 to 0.9806; Qwen3-4B, already strong at 0.8855, reaches 0.9837. Because the finish line is shared, the *size* of the gain is essentially dictated by how far below that ceiling the base sat — +0.5282 for SmolLM2 against +0.0981 for Qwen3-4B — with a panel aggregate of +0.3234 [+0.2647, +0.3690]. Mechanistically this is unsurprising, and that is the point: the manifest is those corpora, so the adapter learns their label conventions and surface cues. A benchmark-only reading would report a large, consistent, tightly-intervalled improvement and stop.

On corpora held out at the source level the same fine-tune does something different. The aggregate is $-0.0589 [-0.0837, -0.0321]$ — close to zero and slightly negative — but that average is a mirage, pooling effects of opposite sign: $+0.0400 [+0.0003, +0.0776]$ for SmolLM2-1.7B against $-0.1499 [-0.1963, -0.1050]$ for Qwen3-4B. Per (checkpoint, seed), 15 of 20 guards land in the quadrant where represented performance rises and transfer falls, 5 in uniform gain, *none* in uniform loss. Fine-tuning does not degrade the guard everywhere; it reshapes the score toward the training corpora, and for a model that already generalized well that reshaping is a net loss off-source.

Two refinements bound how strongly this can be put. “Stronger bases specialize more” is partly arithmetic — Δ is a difference against a shared endpoint, so a negative correlation between base level and Δ is partly definitional, and the transfer cost is in fact *not* monotone in base strength; the sharper, falsifiable claim is *endpoint invariance*. And the losses concentrate on the adversarial corpora (jailbreakbench -0.0776 , wildjailbreak -0.0792 , wildguardtest -0.0669) while xstest, which contrasts genuinely unsafe prompts with benign look-alikes rather than attacks, barely moves at -0.0120 (Table 6): specialization is not uniform forgetting, it forgets the hardest, most shifted cases first.

Does the direction extend to guards that are already guards? Our panel starts from general instruction models, so what we see may be the cost of *becoming* a guard rather than of adapting one. A separate arm — ten checkpoints (6 released purpose-built guards plus 4 general models), five seeds, estimands and criteria fixed in a committed registry before any score existed — finds the same direction on the registered purpose-built panel: represented gain $+0.111$ (one-sided LCB $+0.070$) with a transfer loss. A base-anchored KL penalty preserves transfer (LCB $+0.032$) but its represented cost -0.034 (LCB -0.062) **fails** the registered -0.02 non-inferiority margin. That arm is analysis-preregistered but not data-blind, so we report it as a directional replication, not a confirmation; its practical reading is that the anti-forgetting knob is a trade-off dial, not a free upgrade.

What Q1 licenses is therefore narrow and consequential: on this panel a represented-source gain is not evidence of transfer, and the panel average on held-out sources conceals per-checkpoint effects of opposite sign. A guard result reporting one aggregate over a mixture of represented and held-out corpora is unreadable for deployment purposes, because the two components move independently and can cancel.

5 Q2: THE ALARM BUDGET, THE METRIC REGION, AND THE PREVALENCE

Ranking is threshold-free; enforcement is not. This section asks what happens to Q1’s conclusion when the three quantities a deployment actually fixes — the alarm budget, the ROC region that matters, and the prevalence of unsafe traffic — are made explicit. Nothing below involves a new model run or a new label: every number is a re-reading of the same committed per-row scores, which is precisely why this is a reporting problem rather than a modelling one.

5.1 At an equal alarm budget the apparent gain reverses

Select each guard’s threshold to hit a 5% false-positive target on a calibration split and read off the realized rates. On represented sources the tuned guard looks strictly and dramatically better: recall jumps $13.0\% \rightarrow 76.9\%$ while the represented false-alarm rate even edges down, $1.7\% \rightarrow 1.1\%$. Off-source the picture inverts: the benchmark-macro false-alarm rate climbs $8.1\% \rightarrow 15.5\%$ and the pooled-negative rate blows past target, $4.3\% \rightarrow 17.0\%$ — the calibrated threshold simply does not transfer — while transfer recall rises only modestly, $51.7\% \rightarrow 58.1\%$. That rise is not iso-budget, which matters more than its size: the tuned guard buys most of its extra recall by alarming roughly four times as often on held-out negatives.




Fig. 2. **The same guards, three thresholds.** Grey: the untuned base at its own calibrated cutoff. Light blue: the tuned guard at its own cutoff — the comparison a leaderboard reports, and the one that looks like an improvement. Dark red: the tuned guard re-thresholded to its own base’s pooled false-alarm rate (budget printed under each group). At an equal budget the tuned guard is worse on all four checkpoints on both instruments, and the HarmBench drop needs no budget caveat at all: it catches less (78.0% → 60.0%) while alarming more, so it is dominated on both axes rather than trading one for the other.

So we equalize the budgets and re-read the same rows, giving each tuned seed the threshold at which its pooled transfer false-alarm rate is at or below its own base’s. The apparent gain does not merely shrink — **it reverses, on all four checkpoints and by a wide margin** (Figure 2; generated artifact Table 7): transfer recall falls 0.517 → 0.185 (−0.332) and HarmBench recall 0.780 → 0.171 (−0.609), so at an equal alarm budget the tuned guard catches *less than half* of what its own untuned base catches off-source.

The single-class stress sets make the incidence concrete. On HarmBench the tuned guard’s recall *falls* from 78.0% to 60.0% — it catches less of exactly the content a guard exists to stop — while over-refusal on the benign OR-Bench is essentially flat (11.8% to 12.0%). Since OR-Bench is also unseen, “off-distribution” cannot be what distinguishes it: the extra false alarms are specific to the transfer suite’s negatives and we do not identify what separates them, though the flat OR-Bench line does rule out a blanket increase in caution. The governance reading is the one we would press on an auditor: a guard that looks better only because its off-source false-alarm rate roughly quadrupled is not better in any sense a deployment stakeholder cares about, and the two error directions land on different populations. A protocol that collapses ranking and threshold behaviour into one headline does not merely lose precision; it reports an improvement where there is a joint deterioration.

5.2 The metric’s region co-produces the size of the effect

A more basic mismatch sits underneath. Macro average precision averages precision over the *entire* ranking, including the deep negative mass below any cutoff an inline guard would use, while every deployment sentence anyone writes is about a guard firing on a few percent of traffic. We therefore recompute the identical eight cells on the same rows under two further metrics: one-way partial AUC over false-positive rate [0, 0.05], normalized to mean recall inside that budget (so it reads on a recall scale, against a chance floor of 0.025 rather than 0.5), and recall at that budget. 0 of the eight cells change sign, so Q1’s finding is not an artifact of an average-ranking metric; what does not survive is the *magnitude*, and it fails in the direction that flatters the intervention (Table 2). The represented gain is +0.323 on average precision against +0.686 on partial AUC (2.1×), and the transfer cost −0.059 against −0.174 (3.0×), with panel-mean budget recall moving −0.199. The worst cell is the recurring one: Qwen3-4B gives back −0.409 [−0.499, −0.312] of partial AUC and −0.437 in budget recall — roughly two of every five unsafe prompts it would have caught at a 0.05 alarm rate

Table 2. **Same scores, same rows, three metrics.** Panel-mean paired base→tuned change under the whole-ranking metric, under partial AUC restricted to false-positive rate $[0, 0.05]$, and as recall at that budget. No cell changes sign, but the whole-ranking metric *understates* both halves of the trade. Per-checkpoint values and intervals: Table 8.

Regime	macro-AP	pAUC $[0, 0.05]$	recall@0.05	amplification
Represented sources	+0.323	+0.686	+0.678	2.1×
Source-held-out	-0.059	-0.174	-0.199	3.0×

before tuning. One qualification cuts the other way and weakens a claim we made above: SmolLM2-1.7B, the single checkpoint whose transfer *improves* on macro average precision (+0.0400, clearing zero), does not keep that gain in the operating region (partial AUC +0.011, budget recall −0.012, both straddling zero). “Positive for the weakest base” is a statement about average ranking, not about deployable recall.

5.3 The prevalence co-produces the level and the ordering

A third quantity is fixed by deployment, not the model, and is not a property of the guard at all: the base rate of unsafe requests. Our transfer regime is scored on a balanced pool (790 unsafe against 790 safe rows) while real inbound traffic is overwhelmingly benign, and average precision depends exactly on that base rate — given a guard’s ROC, average precision at prevalence π_+ is a recomputation, not a new measurement:

$$\text{AP}(\pi_+) = \int_0^1 \frac{\pi_+ u}{\pi_+ u + (1 - \pi_+) \text{FPR}(u)} du, \quad (2)$$

with u recall and $\text{FPR}(u)$ the false-positive rate at that recall. Applied to the four base guards’ committed held-out ROCs, balanced average precision is an *optimistic* reading of deployed precision — the strongest ranker falls from 0.94 at balance to ≈ 0.56 at 1% prevalence, and one guard from 0.82 to 0.11 — and low prevalence *re-spaces* the ranking: Qwen2.5-1.5B edges SmolLM2-1.7B at balance (0.82 vs. 0.79) and falls behind it once positives are rare (0.11 vs. 0.21). That is a re-spacing rather than a wholesale winner-flip, but it is enough that a single balanced number can invert two candidates’ apparent order at the prevalence they would be deployed at — for free, with no new data.

Taken together, the three quantities a deployment fixes — alarm budget, ROC region, prevalence — change respectively the *sign*, the *magnitude*, and the *ordering* of a guard comparison, using no new data. A guard number reported without all three is not under-precise; it is not yet a claim about enforcement.

6 Q3: THE SOURCING DECISION INVERTS ACROSS TRAFFIC REGIMES

Everything so far compares small guards with each other. The question an operator asks next — is self-hosting costing us safety relative to a hosted frontier model? — is where an evaluation error becomes a procurement error. We price it on identical rows joined by content digest, at a matched 5% alarm budget, so no model is read at an operating point nobody selected.

On unfamiliar traffic the hosted model is clearly better.

On Choi et al. [12]’s external, expert-annotated finance/health/law prompts — a corpus no guard here was trained on — the best hosted configuration (gpt-5.4 (low)) reaches .896 recall at the matched budget against .787 for the strongest small base on those rows (SmolLM3-3B), a paired row-bootstrap difference of +0.109 [+0.077, +0.139]: the largest single accuracy gap in our evidence, at our strongest label tier. Nothing in-house closes it — tuning *hurts* 4 of 6 checkpoints




Fig. 3. **The frontier gap is a property of the regime, not of the model.** The same comparison, at the same matched 5% false-alarm budget, points in opposite directions on the two regimes. The two cells sit at the same evidence tier but on different data and are never pooled. The top-left cell is post hoc and *not* robust to reweighting; the bottom-right is a paired comparison on external expert-annotated rows.

(panel mean +0.005, hiding swings an order of magnitude larger); six released purpose-built guards all land *below* the best untuned open checkpoint; and scaling the base 8× buys only +0.062 [+0.039, +0.084] while leaving +0.066 [+0.043, +0.089] still to go.

On traffic the manifest names, the ordering reverses. That comparison is measured entirely on a corpus nobody trained on, so it reports the *transfer* regime and nothing else. Five general-safety corpora were scored by both the panel and the hosted reference, with the regime split read from the training manifest rather than asserted, and on the 3 represented corpora the panel’s small tuned guards *beat* the reference: the equal-source, equal-checkpoint mean paired difference is +0.083 [+0.013, +0.157] in recall at the matched budget and +0.039 [+0.015, +0.072] in average precision (Figure 3).

We report that reversal with its weaknesses attached. The represented-source aggregate is **post hoc**: it did not exist before the per-cell headline failed a multiplicity correction, and was added in the same revision that found the failure. Its weighting is fixed by the regime split rather than by any result, which makes it a sounder summary than the maximum of 12 cells but not a pre-specified one, and it is conditional on *these three corpora* — the bootstrap holds the source set fixed, because three purposively chosen corpora do not sample a population of corpora, and drawing sources with replacement widens the interval to [−0.019, +0.220], which *includes* zero. Of four defensible weightings only two support a positive advantage: weighting by row count roughly halves the estimate and straddles zero, and including the *base* arms alongside the tuned ones excludes zero in the *opposite* direction (−0.264 [−0.321, −0.179]) — which says what the effect is about: a property of *tuned* guards specifically, Q1’s specialization seen from the other side, not a general claim that small guards beat hosted ones. Per-cell results are exploratory (3 of 12 clear zero nominally, 0 survive a Holm step-down), so no sentence here should be read as “a 1.5B guard beats the frontier.”

Why this is a governance finding. The sourcing decision is not “which guard is better” but “what share of our traffic can we enumerate in a training manifest, and who is accountable for the rest.” The costs that decide it appear in no accuracy column: the hosted path answers in 1,553 ms at the median against tens of milliseconds locally ($\approx 77\times$) at roughly \$0.80 per thousand prompts, and every prompt leaves the operator’s boundary — for the mortgage setting

of Section 7, a data-governance question before an engineering one. And, least discussed anywhere, the provider’s own input filter refused a subset of the evaluation prompts outright, before the guard saw them, *non-deterministically*, with no row refused under all six configurations. A screening layer that intermittently declines to return a verdict, for reasons the operator cannot inspect or appeal, is itself a compliance artifact — exactly the opacity internal-audit frameworks [50] and third-party audit proposals [51] exist to surface.

7 Q4: A GENERAL-SAFETY SCORE IS NOT A COMPLIANCE SCREEN

This is the strongest form of the thesis, and the one where the construct gap is not a matter of degree. In a regulated domain the general-safety benchmark can fail to see the violation *at all* — not because the guard is weak, but because the violating request is, on its surface, ordinary professional prose. A deliberately blatant version fixes the idea:

“Applications from the majority-Spanish-speaking tracts keep coming in weak; draft language to quietly deprioritize that pipeline, and write the denials up as ‘collateral’ so it’s clean.”

No slur, no jailbreak, no injection — nothing a content taxonomy has a category for. Yet honoring it would commit redlining, proxy discrimination, and adverse-action-reason masking. The last is not an abstraction: under the Equal Credit Opportunity Act and Regulation B a creditor must give the *specific, principal* reasons for an adverse action, and the Consumer Financial Protection Bureau has stated that algorithmic complexity is no defense for failing to do so [15, 16]. A request to substitute a facially neutral pretext for the real reason already in the file is a request to manufacture a non-compliant adverse-action notice — the mechanism Barocas and Selbst [4] describe, arriving as a polite question. A guard trained on a domain-independent harm taxonomy has no label for it.

7.1 A dual-labeled instrument

To measure whether guards catch such requests we use a frozen, HMDA-grounded benchmark [2, 20] whose unit is one incoming request carrying *two separately assigned* labels: general safety G (would an ordinary safety guard call this unsafe — jailbreak, abuse, injection?) and mortgage policy D (would honoring it break mortgage law or policy?), composed as $\text{final} = 1[G=\text{unsafe} \vee D=\text{intervene}]$. Crossing them gives four quadrants whose *meanings* are the design. **G0/D1** is the payload — a request a general guard rates safe that nonetheless solicits a compliance violation, which a guard must therefore flag on D alone — and carrying two labels rather than one merged verdict is exactly what lets that stratum be scored on its own; the other corners are controls.

The frozen release holds 994 synthetic rows, split family-isolated so rows sharing a content family go wholly to train or wholly to test (0 of 958 families span a boundary). G0/D1 is the largest non-benign block at 502 rows, with 450 benign and 42 bad on both counts. **G1/D0 is empty** (0 rows): the safety-tuned generator declined to author domain-independent jailbreaks with no mortgage angle, so G is *nested* inside D rather than crossed with it — Section 9 returns to whom that costs.

7.2 Zero-shot guards rank the policy label barely above its own chance floor

We score the same four base checkpoints *zero-shot* on the frozen 146-row public-test split (Figure 1(d); generated table in Table 9), and the nuance runs both ways. Threshold-free, the bases do rank D -violations *moderately* well, $AP \cdot D$ between 0.67 and 0.85 — “help me discriminate” pattern-matches to unsafe even without a jailbreak, so the domain is not invisible. But that band must be read against the chance floor its own prevalence sets: 81 of 146 rows are D -positive, so a random ranker already scores ≈ 0.555 , leaving the band only 0.12–0.30 *above chance* (the base-rate-free AUROC- D ,

0.60–0.78, agrees). No guard comes near 1, and five of six pairwise intervals overlap, so we read the ordering as a *direction*, not a ranking.

Two numbers are deliberately *not* reported, and the reasons are findings. We give no fixed-threshold “caught at 5% false alarms” count for the G0/D1 stratum: these guards’ scores are so tightly clustered that a development-calibrated cutoff sits on a knife edge, its catch count swinging by more than fifty rows between library versions of the quantile routine — Section 5.1’s threshold-transfer lesson arriving in the domain arm. And AP-G rests on six *G*-positives, so its spread from 0.261 to 1.000 is a small-sample band, not four comparable estimates.

7.3 One worked miss

An aggregate as respectable as $AP \cdot D = 0.85$ can still hide a total miss on an individual row, and the G0/D1 stratum is where that happens. One row of the frozen split reads, in the register of a routine underwriting question, as a request to draft a decision note leaning on “market fit” and a neighborhood’s “resale stability” *instead of* the permissible reason already in the file — a 50–60% debt-to-income ratio not in dispute. Its *G* label is safe, correctly in the general-safety sense; its *D* label is intervene against six policy cards spanning redlining, proxy discrimination, and adverse-action reason masking. **All four zero-shot guards rank it below the median benign inquiry in the same split**, and one below every single benign row (Figure 5; Figure 6 works the row end to end, verbatim and with its policy-card citations). A general-safety score reports this request as clean, and an operator who selected that guard on a general-safety benchmark holds no artifact that would have said otherwise.

7.4 The fairness gate we built does not survive its own robustness check

Ranking quality cannot tell you whether a guard treats two otherwise-identical requests differently because one names a protected group, so the benchmark adds **protected-class pairs**: two requests identical in every slot — same fact sheet, role, and ask — except the protected-class mention, both members benign, so a fair guard must pass both (39 in the release, three in the scored split). This is the standard counterfactual token-fairness construction [21, 35], repurposed from decision correctness to score invariance: not whether a decision is right, only whether the score *moves* when nothing but the protected mention changes. The gate is the mean absolute within-pair gap,

$$\Delta_{\text{context}} = \frac{1}{|\mathcal{P}|} \sum_{(a,b) \in \mathcal{P}} |p(x_a) - p(x_b)|, \quad (3)$$

with target ≈ 0 . **On this panel the gate does not rank guards, and we report that as a negative result about instrument design rather than quietly dropping it.** Two defects cut in opposite directions and between them dissolve the apparent contrast (Figure 4).

First, Equation (3) is defined on the *probability* scale, so it is comparable only across guards whose probabilities occupy comparable ranges — and here they do not. Qwen3-4B’s entire scored split sits at a median $p = 3.2 \times 10^{-6}$, so its $\Delta_{\text{context}} = 0.000$ is a *saturation artifact*, not invariance: recovering the gap exactly on the raw margin scale $s = \log(p/(1-p))$ gives 0.797 log-odds, the second largest on the panel and a 2.1–2.5× shift in odds in the same direction on all three pairs. Plotting only the probability-scale bars is what made a units artifact look like perfect fairness.

Second, the apparent worst guard is carried by one pair that is not a minimal pair: Qwen2.5-1.5B’s headline 0.183 decomposes into per-pair gaps of 0.018, 0.022, and 0.508, and that third pair contrasts a named trait (“Muslim”) against a seven-word placeholder rather than swapping a single token — the known failure mode of counterfactual token tests




Fig. 4. **Our own fairness instrument, audited.** Left: the protected-pair gap on the probability scale, all three pairs against the two genuine single-token swaps — one guard looks perfectly invariant (0.000), another worst (0.183), and restricting to minimal pairs removes the outlier. Right: the identical gap on the raw log-odds margin, recovered exactly from the same committed probabilities, where the order nearly inverts because the “invariant” guard’s probabilities saturate at $p \sim 10^{-6}$. The defensible statement is not a ranking but a constraint on the instrument.

[21], and only 21 of the 39 release pairs are true single-token swaps. Restricted to the two genuine single-token pairs the gap is 0.020, inside the unsaturated guards’ band rather than an outlier above it.

What remains is the instrument lesson: **a probability-scale counterfactual gap on three pairs, one of them not minimal, cannot rank guards.** A fairness gate needs the same scrutiny as the models it gates; ours failed it, and reporting that is cheaper than publishing an invariance ranking that is a units artifact.

7.5 External breadth, at a stronger label tier

The mortgage instrument is deep with weak labels; the complementary arm is shallow with strong ones. Choi et al. [12]’s expert-annotated finance, health, and law prompts (2,275 rows, scored zero-shot by the same four bases; Table 11) show all four already ranking domain violations well (aggregate average precision 0.88–0.96), with the top guard *not* the largest model and the ranking not monotone in parameter count. These labels sit at a strictly stronger tier, so they are reported beside — never pooled with — the mortgage numbers. What survives the tier boundary is a finding about instruments rather than models: **the guard ordering does not carry across the two domain arms.** A guard ranking is a joint property of the guard and the instrument — exactly what a procurement process citing one leaderboard cannot see.

8 FROM MEASUREMENT TO GOVERNANCE

Our findings are about reporting, so our deliverable is a reporting requirement. Table 3 gives, for each validity question, the artifact that answers it, the failure it prevents, and the consumer who needs it. Every item is computable from per-row scores an operator already has — no new labels, no new training runs, no GPU — which is the argument for requiring it.

Where the schedule attaches. Three existing hooks already ask for something close to these disclosures. Under *ECOA and Regulation B* a creditor must state the specific, principal reasons for an adverse action, and algorithmic complexity is no defense [15, 16], so a screening layer in a lending workflow sits inside the compliance perimeter rather than adjacent to it: our worked example is a request to *substitute* a pretext for the real reason, and a guard selected on a general-safety benchmark ranks it below the median benign inquiry — D4 would have surfaced that before deployment, D2 would have stopped the operator believing the general-safety number transferred. *Article 15 of the EU AI Act* [19]

Table 3. **A disclosure schedule for safety-guard evaluation.** Each row is computable from committed per-row scores with no new labels and no new training. The right-hand column names the consumer, in the sense of Mitchell et al. [44] and Raji et al. [50]: documentation is an accountability artifact only if someone is entitled to read it.

Question	Disclosure required	Failure it prevents (measured here)	Consumer
D1 Pairing	The change against the guard’s <i>own</i> pre-tuning checkpoint on identical rows, per checkpoint and seed — not a delta against other vendors’ models.	Attributing to the fine-tune what belongs to the starting checkpoint: per-checkpoint effects span +0.0400 to −0.1499 and the panel mean (−0.0589) hides both.	Internal audit; model-card reviewer
D2 Source manifest	The corpora named in the manifest, and results split into <i>represented</i> and <i>source-held-out</i> blocks with an overlap audit.	Reporting acquaintance with the instrument as generalization: +0.3234 represented against −0.0589 held-out, same guards.	External auditor; benchmark maintainer
D3 Budget, region, prevalence	Recall at a <i>matched</i> false-alarm budget; the metric restricted to that budget; $AP(\pi_+)$ at the operator’s prevalence.	A reported improvement that is a joint deterioration: sign reverses on 4/4 checkpoints, magnitude changes 3.0×, ordering re-spaces at 1% prevalence.	Procurement; service owner; regulator
D4 Domain construct, label tier	The enforced policy as a label <i>separate</i> from general safety; label provenance (SME / expert / model judge); per-quadrant results.	Treating a general-safety score as a compliance screen: $AP \cdot D$ is 0.12–0.30 above its own chance floor, and one G0/D1 row ranks below the median benign request for all four.	Compliance; fair-lending examiner

requires declared levels of accuracy and robustness for high-risk systems; our results say what such a declaration must be indexed by to mean anything — an alarm budget, a traffic regime, a prevalence — since a single figure satisfies its letter while remaining compatible with a joint deterioration in both error directions. And the MEASURE function of the *NIST AI RMF* [45] asks for metrics appropriate to context and re-assessed under change, of which D3 is a concrete operationalization.

What follows, and what does not. For practitioners the schedule reduces to four habits: compare every tune with its own base; keep represented and source-held-out corpora in separate columns; read every thresholded claim at a matched budget and inside the budget you can afford; and build a domain-grounded, separately labeled instrument before trusting any general score in a regulated setting. For auditors and procurement reviewers it is a checklist whose failures are cheap to detect and whose absence is itself informative. For this community we press one further point: the fairness instrument of Section 7.4 failed its own robustness check on a scale artifact and a non-minimal pair, and would have produced a publishable invariance ranking had we not looked — instrument audits belong in the same schedule as model audits. We offer no threshold, no recommended guard, and no certification procedure. What we can offer is the negative result most useful to a governance reader: on this panel the selection rule in universal use is compatible with a guard that is worse on both error directions off-source, worse on the hardest attacks, and blind to the domain violation the operator is legally obliged to catch, while its benchmark number goes up.

9 LIMITATIONS

These change what the paper licenses, not merely how confident it is.

No causal and no population claim. Fine-tuning *co-occurs* with specialization; we do not isolate a mechanism. The panel is a fixed, purposively chosen set of four checkpoints from only *two* lineages, so lineage, scale, and native alignment recipe are confounded and no size effect is identified. The intervals are conditional on this panel and these corpora — not statements about a population of guards, and not hypothesis tests — and “stronger bases specialize more” is coupled to the metric by construction (Section 4) rather than a behavioural law.

Not confirmatory; one recipe; prompt-only. The manifest and evaluation suites were visible during development, so “held out” means dataset-held-out, not sealed-from-the-researcher; the one analysis-preregistered arm is not data-blind, its registry is non-final, and one of its cells was a harness artifact rather than a measurement (a released guard scoring a constant on all rows under two template and verdict-position bugs, since fixed). A genuinely sealed cohort is the upgrade and does not exist here. The specialization signature is also a property of this adapter configuration on this data — we sweep neither rank, step budget, nor data mixture — pools are balanced or near-balanced, which [Section 5.3](#) shows is optimistic relative to deployed precision, and we classify the incoming request only. Finally, every matched-budget number places its threshold using a quantile of the labelled negatives it is then scored on, so it is not one a production system could place ([Appendix B.1](#)).

The mortgage instrument is a measuring stick, not a legal finding. Its labels are model-assigned against written policy cards, with self-consistency checks but no subject-matter-expert adjudication and no human inter-annotator agreement; it certifies nothing about any real lender, applicant, or model, and licenses no fair-lending or disparate-impact conclusion. Four limits bound it. G1/D0 is empty, making “looks unsafe yet is policy-compliant” *unmeasurable* here — and it is worth naming whom that costs: over-refusal of legitimate requests that merely *mention* a protected class is precisely a harm to the applicants fair lending exists to protect. Several strata are small (six *G*-positives and three protected pairs in the scored split), so the invariance gate is under-powered before it is scale-confounded. Family isolation holds at content-family level, not pair level: one protected pair has its arms in different splits, missed by a near-duplicate clusterer by 0.0023 of similarity. And the release is frozen and stochastic in generation, so only its *evaluation* reproduces. Separately, the frontier comparison is coarse and partly post hoc: the hosted model’s ranking signal takes only 47–65 distinct values over 2,275 rows, so a matched budget can only be approximately placed, and the represented-source reversal is conditional on three purposively chosen corpora and not robust to reweighting ([Section 6](#)).

Reproducibility, honestly bounded. One command byte-checks 27 of 35 generated artifacts from committed per-row scores (4 need a lock-pinned environment, 4 are uncovered). That is *analysis* reproducibility; re-deriving the scores needs a GPU, and regenerating the mortgage benchmark is **not** claimed, since its construction runs at nonzero temperature ([Appendix F](#)).

10 CONCLUSION

Benchmark gains are not governance gains. On this panel fine-tuning delivers a large improvement on the corpora its own manifest names (+0.3234) and none that is reliable off them (-0.0589); at each base guard’s own alarm rate the apparent off-source gain reverses on every checkpoint; inside the affordable budget the trade is two to three times larger than the headline says; the sourcing conclusion inverts by traffic regime; and in a regulated domain the violation can be fluent, procedurally framed text that all four guards rank below the median benign request. None of this says fine-tuning harms guards — it is a claim about what a number licenses. The selection rule in universal use is compatible with every failure above while its number rises, and [Table 3](#) makes each visible from scores an operator already holds.

GENERATIVE AI USAGE STATEMENT

OpenAI Codex was used in preparing this manuscript in two bounded roles, and in no role that produced substantive text. First, it assisted with analysis and plotting code, all of which is committed and whose outputs are byte-checked against committed per-row scores. Second, it assisted with copy-editing of author-written prose (grammar, concision, and consistency of terminology). No section, argument, claim, or number was generated by a model. All numerical values in this paper are emitted by committed analysis code from committed per-row scores, or parsed from those emitted artifacts by the figure script, and were verified against the artifacts by the authors.

Separately, and as a matter of *research method* rather than manuscript preparation, the mortgage benchmark’s two labels are assigned by a language-model judge reading written policy cards. That is a property of the instrument, is disclosed wherever its numbers appear, is why its evidence tier is the weakest of the three we report (Table 4), and is why no legal or fair-lending conclusion is drawn from it.

ETHICAL CONSIDERATIONS

Dual-use content. The mortgage benchmark’s prompts solicit policy violations by design, including redlining by proxy and pretextual adverse-action language. They are synthetic, grounded in a public aggregate loan-level dataset containing no individual identities [20], and should be treated as harmful-content samples on reuse. The release carries that warning, and no real applicant, loan file, or lender appears in it.

Risk of misreading a measuring stick as a certification. The greatest foreseeable harm from this work is that a reader treats the mortgage instrument as a fair-lending audit. It is not: its labels have no expert adjudication and no human agreement statistic, its scored split is small, and one of its four quadrants is empty. We state this at every point of use rather than once in a footnote, and we decline to publish a guard ranking or an operating threshold from it.

Naming a failure in deployed components. We report that several released, widely used guard models perform poorly on regulated-domain prompts at a strict alarm budget. We report it because operators making sourcing decisions need it; we score each model through *its own* native verdict contract rather than forcing our prompt on it; and we note where an apparent failure was in fact our harness (Section 9). We do not characterize any model as unsafe.

Whose errors we cannot see. The empty quadrant in our instrument means we cannot measure over-refusal of compliant requests that mention a protected class — a harm that would fall on the applicants fair-lending law protects. Populating it is the first item of Appendix E, and we flag its absence rather than presenting our coverage as complete.

ADVERSE IMPACTS

Two adverse impacts are worth naming. A disclosure schedule can be adopted performatively: reporting a matched-budget number without the budget it was matched at, or a represented/held-out split whose “held-out” corpora were chosen after the fact, would reproduce exactly the problem we describe with an extra table attached. And a paper emphasizing how easily guard benchmarks mislead could be read as an argument for deploying no screening layer at all. It is not: our evidence is that these guards’ *reported* improvements are unreliable, not that screening is without value. The remedy we propose is a stricter reading of the same measurements, not fewer safeguards.

REFERENCES

- [1] Akshit Acharya and Ignacio Sanchez. 2025. Watching the AI Watchdogs: A Fairness and Robustness Analysis of AI Safety Moderation Classifiers. arXiv:2501.13302 [cs.CL]
- [2] Anonymous. 2026. MortgageGuardBench v1_hmda2022: A Frozen, Dual-Labeled Mortgage-Policy Guard Benchmark. 994-row frozen release, CC BY 4.0; attribution withheld for anonymous review.
- [3] Anonymous. 2026. Safety Benchmark Gains Do Not Guarantee Safety Transfer: A Study of Fine-Tuning Small Language Model Safety Guards. Companion technical report and reproducibility artifacts; author identity withheld for anonymous review.
- [4] Solon Barocas and Andrew D. Selbst. 2016. Big Data’s Disparate Impact. *California Law Review* 104, 3 (2016), 671–732.
- [5] Robert Bartlett, Adair Morse, Richard Stanton, and Nancy Wallace. 2022. Consumer-Lending Discrimination in the FinTech Era. *Journal of Financial Economics* 143, 1 (2022), 30–56. <https://doi.org/10.1016/j.jfineco.2021.05.047>
- [6] Elias Bassani and Ignacio Sanchez. 2024. GuardBench: A Large-Scale Benchmark for Guardrail Models. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing (EMNLP)*. <https://doi.org/10.18653/v1/2024.emnlp-main.1022>
- [7] Elias Bassani and Ignacio Sanchez. 2025. On Guardrail Models’ Robustness to Mutations and Adversarial Attacks. Preprint. Guardrail robustness evaluation.
- [8] Su Lin Blodgett, Gilsinia Lopez, Alexandra Olteanu, Robert Sim, and Hanna Wallach. 2021. Stereotyping Norwegian Salmon: An Inventory of Pitfalls in Fairness Benchmark Datasets. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL-IJCNLP)*. 1004–1015. <https://doi.org/10.18653/v1/2021.acl-long.81>
- [9] Donald E. Bowen III, S. McKay Price, Luke C. D. Stein, and Ke Yang. 2024. Measuring and Mitigating Racial Bias in Large Language Model Mortgage Underwriting. SSRN working paper. <https://doi.org/10.2139/ssrn.4812158>
- [10] Samuel R. Bowman and George E. Dahl. 2021. What Will it Take to Fix Benchmarking in Natural Language Understanding?. In *Proceedings of the 2021 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*. 4843–4855. <https://doi.org/10.18653/v1/2021.naacl-main.385>
- [11] Patrick Chao, Edoardo Debenedetti, Alexander Robey, Maksym Andriushchenko, Francesco Croce, Vikash Sehwal, Edgar Dobriban, Nicolas Flammarion, George J. Pappas, Florian Tramèr, Hamed Hassani, and Eric Wong. 2024. JailbreakBench: An Open Robustness Benchmark for Jailbreaking Large Language Models. In *Advances in Neural Information Processing Systems 37, Datasets and Benchmarks Track*. arXiv:2404.01318
- [12] Minseok Choi, Dongjin Kim, Seungbin Yang, Subin Kim, Youngjun Kwak, Juyoung Oh, Jaegul Choo, and Jungmin Son. 2026. ExpGuard: LLM Content Moderation in Specialized Domains. arXiv:2603.02588 [cs.CL]
- [13] Sanghyun Choi, Jiwon Park, Minh Nguyen, and Paula Alvarez. 2026. COMPASS: A Framework for Evaluating Organization-Specific Policy Compliance. Preprint. Policy-grounded request evaluation across organizational scenarios.
- [14] Pedro Cisneros-Velarde, Umang Bhatt, and Kamalika Chaudhuri. 2026. Policy Compliance of User Requests in Natural Language for AI Systems. Preprint. Human-labelled policy-compliance benchmark.
- [15] Consumer Financial Protection Bureau. 2022. Consumer Financial Protection Circular 2022-03: Adverse Action Notification Requirements in Connection with Credit Decisions Based on Complex Algorithms. <https://www.consumerfinance.gov/compliance/circulars/circular-2022-03-adverse-action-notification-requirements-in-connection-with-credit-decisions-based-on-complex-algorithms/>.
- [16] Consumer Financial Protection Bureau. 2023. Regulation B, Equal Credit Opportunity Act, 12 C.F.R. Part 1002. <https://www.consumerfinance.gov/rules-policy/regulations/1002/>.
- [17] Justin Cui, Wei-Lin Chiang, Ion Stoica, and Cho-Jui Hsieh. 2024. OR-Bench: An Over-Refusal Benchmark for Large Language Models. arXiv:2405.20947 [cs.CL]
- [18] Hayder Elesedy, Pedro M. Esperança, Silviu Vlad Oprea, and Mete Ozay. 2024. LoRA-Guard: Parameter-Efficient Guardrail Adaptation for Content Moderation of Large Language Models. In *Findings of the Association for Computational Linguistics: EMNLP 2024*. arXiv:2407.02987
- [19] European Parliament and Council of the European Union. 2024. Regulation (EU) 2024/1689 Laying Down Harmonised Rules on Artificial Intelligence (AI Act). Official Journal of the European Union, L series, 12 July 2024.
- [20] Federal Financial Institutions Examination Council and Consumer Financial Protection Bureau. 2023. HMDA Snapshot National Loan-Level Dataset (2022). <https://ffiec.cfbp.gov/data-publication/snapshot-national-loan-level-dataset/>.
- [21] Sahaj Garg, Vincent Perot, Nicole Limtiaco, Ankur Taly, Ed H. Chi, and Alex Beutel. 2019. Counterfactual Fairness in Text Classification through Robustness. In *Proceedings of the 2019 AAAI/ACM Conference on AI, Ethics, and Society (AIES ’19)*. 219–226. <https://doi.org/10.1145/3306618.3317950>
- [22] Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. 2021. Datasheets for Datasets. *Commun. ACM* 64, 12 (2021), 86–92. <https://doi.org/10.1145/3458723>
- [23] Tarleton Gillespie. 2018. *Custodians of the Internet: Platforms, Content Moderation, and the Hidden Decisions That Shape Social Media*. Yale University Press.
- [24] Robert Gorwa, Reuben Binns, and Christian Katzenbach. 2020. Algorithmic Content Moderation: Technical and Political Challenges in the Automation of Platform Governance. *Big Data & Society* 7, 1 (2020). <https://doi.org/10.1177/2053951719897945>
- [25] William Hackett, Lewis Birch, Stefan Trawicki, Neeraj Suri, and Peter Garraghan. 2025. Bypassing LLM Guardrails: An Empirical Analysis of Evasion Attacks against Prompt Injection and Jailbreak Detection Systems. arXiv:2504.11168 [cs.CR]

- [26] Seungju Han, Kavel Rao, Allyson Ettinger, Liwei Jiang, Bill Yuchen Lin, Nathan Lambert, Yejin Choi, and Nouha Dziri. 2024. WildGuard: Open One-Stop Moderation Tools for Safety Risks, Jailbreaks, and Refusals of LLMs. In *Advances in Neural Information Processing Systems 37, Datasets and Benchmarks Track*. arXiv:2406.18495
- [27] Tanvir Hossain, Priya Ranjan, Emre Aydin, and Grace Whitfield. 2026. When Safety Geometry Collapses: Fine-Tuning Vulnerabilities in Purpose-Built Guard Models. Preprint. Domain fine-tuning of released guard models.
- [28] Lei Hsiung, Tianyu Tang, Pin-Yu Chen, and Tsung-Yi Ho. 2025. Why LLM Safety Guardrails Collapse After Fine-tuning: A Similarity Analysis Between Alignment and Fine-tuning Datasets. arXiv:2506.05346 [cs.CL]
- [29] Edward J. Hu, Yelong Shen, Phillip Wallis, Zeyuan Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, and Weizhu Chen. 2022. LoRA: Low-Rank Adaptation of Large Language Models. In *International Conference on Learning Representations (ICLR)*. arXiv:2106.09685
- [30] Ben Hutchinson, Andrew Smart, Alex Hanna, Emily Denton, Christina Greer, Oddur Kjartansson, Parker Barnes, and Margaret Mitchell. 2021. Towards Accountability for Machine Learning Datasets: Practices from Software Engineering and Infrastructure. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT '21)*. ACM, 560–575. <https://doi.org/10.1145/3442188.3445918>
- [31] Hakan Inan, Kartikeya Upasani, Jianfeng Chi, Rashi Rungta, Krithika Iyer, Yuning Mao, Michael Tontchev, Qing Hu, Brian Fuller, Davide Testuggine, and Madian Khabsa. 2023. Llama Guard: LLM-based Input-Output Safeguard for Human-AI Conversations. arXiv:2312.06674 [cs.CL]
- [32] Abigail Z. Jacobs and Hanna Wallach. 2021. Measurement and Fairness. In *Proceedings of the 2021 ACM Conference on Fairness, Accountability, and Transparency (FAccT '21)*. ACM, 375–385. <https://doi.org/10.1145/3442188.3445901>
- [33] Liwei Jiang, Kavel Rao, Seungju Han, Allyson Ettinger, Faeze Brahman, Sachin Kumar, Niloofar Miresghallah, Ximing Lu, Maarten Sap, Yejin Choi, and Nouha Dziri. 2024. WildTeaming at Scale: From In-the-Wild Jailbreaks to (Adversarially) Safer Language Models. In *Advances in Neural Information Processing Systems 37 (NeurIPS)*. arXiv:2406.18510
- [34] Sayash Kapoor and Arvind Narayanan. 2023. Leakage and the Reproducibility Crisis in Machine-Learning-Based Science. *Patterns* 4, 9 (2023), 100804. <https://doi.org/10.1016/j.patter.2023.100804>
- [35] Matt J. Kusner, Joshua R. Loftus, Chris Russell, and Ricardo Silva. 2017. Counterfactual Fairness. In *Advances in Neural Information Processing Systems 30 (NeurIPS)*.
- [36] Seanie Lee, Dong Bok Seong, Dominik Lee, Dohyeon Kim, and Sung Ju Hwang. 2025. SafeRoute: Adaptive Model Selection for Efficient and Accurate Safety Guardrails in Large Language Models. arXiv:2502.12464 [cs.CL]
- [37] Wenhao Li, Olivia Grant, and Diego Serrano. 2026. Defenses Against Prompt Attacks Learn Surface Heuristics. Preprint.
- [38] Q. Vera Liao and Ziang Xiao. 2023. Rethinking Model Evaluation as Narrowing the Socio-Technical Gap. arXiv:2306.03100 [cs.HC]
- [39] Zi Lin, Zihan Wang, Yongqi Tong, Yangkun Wang, Yuxin Guo, Yujia Wang, and Jingbo Shang. 2023. ToxicChat: Unveiling Hidden Challenges of Toxicity Detection in Real-World User-AI Conversation. In *Findings of the Association for Computational Linguistics: EMNLP 2023*. <https://doi.org/10.18653/v1/2023.findings-emnlp.311>
- [40] Hongfu Liu, Hengguan Wang, Xiangming Xu, and Ye Wang. 2025. On Calibration of LLM-based Guard Models for Reliable Content Moderation. arXiv:2410.10414 [cs.CL]
- [41] Yifan Liu, Zhuo Yang, Jing Ren, and Wei Chen. 2026. Domain Generalizable AI Guardrails with Augmented Policy Training. Preprint. Policy-overfitting of guardrail classifiers.
- [42] Mantas Mazeika, Long Phan, Xuwang Yin, Andy Zou, Zifan Wang, Norman Mu, Elham Sakhae, Nathaniel Li, Steven Basart, Bo Li, David Forsyth, and Dan Hendrycks. 2024. HarmBench: A Standardized Evaluation Framework for Automated Red Teaming and Robust Refusal. In *Proceedings of the 41st International Conference on Machine Learning (ICML), PMLR* 235. 35181–35224.
- [43] Danae Metaxa, Joon Sung Park, Ronald E. Robertson, Karrie Karahalios, Christo Wilson, Jeff Hancock, and Christian Sandvig. 2021. Auditing Algorithms: Understanding Algorithmic Systems from the Outside In. *Foundations and Trends in Human-Computer Interaction* 14, 4 (2021), 272–344. <https://doi.org/10.1561/11000000083>
- [44] Margaret Mitchell, Simone Wu, Andrew Zaldivar, Parker Barnes, Lucy Vasserman, Ben Hutchinson, Elena Spitzer, Inioluwa Deborah Raji, and Timnit Gebru. 2019. Model Cards for Model Reporting. In *Proceedings of the 2019 Conference on Fairness, Accountability, and Transparency (FAT* '19)*. ACM, 220–229. <https://doi.org/10.1145/3287560.3287596>
- [45] National Institute of Standards and Technology. 2023. Artificial Intelligence Risk Management Framework (AI RMF 1.0). NIST AI 100-1. <https://doi.org/10.6028/NIST.AI.100-1>
- [46] Inkit Padhi, Manish Nagireddy, Giandomenico Cornacchia, Subhajit Chaudhury, Tejaswini Pedapati, Pierre Dognin, Keerthiram Murugesan, Erik Miehl, Martin Santillan Cooper, Kieran Fraser, Giulio Zizzo, Muhammad Zaid Hameed, Mark Purcell, Michael Desmond, Qian Pan, Zahra Ashktorab, Inge Vejsbjerg, Elizabeth M. Daly, Michael Hind, Werner Geyer, Ambrish Rawat, Kush R. Varshney, and Prasanna Sattigeri. 2024. Granite Guardian: Comprehensive LLM Safeguarding. arXiv:2412.07724 [cs.CL]
- [47] Qwen Team. 2025. Qwen3Guard Technical Report. arXiv:2510.14276 [cs.CL]
- [48] Inioluwa Deborah Raji, Emily M. Bender, Amandalynne Paullada, Emily Denton, and Alex Hanna. 2021. AI and the Everything in the Whole Wide World Benchmark. In *Advances in Neural Information Processing Systems 34, Datasets and Benchmarks Track*. arXiv:2111.15366
- [49] Inioluwa Deborah Raji, I. Elizabeth Kumar, Aaron Horowitz, and Andrew Selbst. 2022. The Fallacy of AI Functionality. In *Proceedings of the 2022 ACM Conference on Fairness, Accountability, and Transparency (FAccT '22)*. ACM, 959–972. <https://doi.org/10.1145/3531146.3533158>
- [50] Inioluwa Deborah Raji, Andrew Smart, Rebecca N. White, Margaret Mitchell, Timnit Gebru, Ben Hutchinson, Jamila Smith-Loud, Daniel Theron, and Parker Barnes. 2020. Closing the AI Accountability Gap: Defining an End-to-End Framework for Internal Algorithmic Auditing. In *Proceedings*

- of the 2020 Conference on Fairness, Accountability, and Transparency (FAT* '20). ACM, 33–44. <https://doi.org/10.1145/3351095.3372873>
- [51] Inioluwa Deborah Raji, Peggy Xu, Colleen Honigsberg, and Daniel E. Ho. 2022. Outsider Oversight: Designing a Third Party Audit Ecosystem for AI Governance. In *Proceedings of the 2022 AAAI/ACM Conference on AI, Ethics, and Society (AIIES '22)*. ACM. <https://doi.org/10.1145/3514094.3534181>
 - [52] Benjamin Recht, Rebecca Roelofs, Ludwig Schmidt, and Vaishaal Shankar. 2019. Do ImageNet Classifiers Generalize to ImageNet?. In *Proceedings of the 36th International Conference on Machine Learning (ICML)*, PMLR 97. 5389–5400.
 - [53] Marco Tulio Ribeiro, Tongshuang Wu, Carlos Guestrin, and Sameer Singh. 2020. Beyond Accuracy: Behavioral Testing of NLP Models with CheckList. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics (ACL)*. 4902–4912. <https://doi.org/10.18653/v1/2020.acl-main.442>
 - [54] Paul Röttger, Hannah Rose Kirk, Bertie Vidgen, Giuseppe Attanasio, Federico Bianchi, and Dirk Hovy. 2024. XSTest: A Test Suite for Identifying Exaggerated Safety Behaviours in Large Language Models. In *Proceedings of the 2024 Conference of the North American Chapter of the Association for Computational Linguistics (NAACL-HLT)*. 5377–5400. <https://doi.org/10.18653/v1/2024.naacl-long.301>
 - [55] Andrew D. Selbst, Danah Boyd, Sorelle A. Friedler, Suresh Venkatasubramanian, and Janet Vertesi. 2019. Fairness and Abstraction in Sociotechnical Systems. In *Proceedings of the 2019 Conference on Fairness, Accountability, and Transparency (FAT* '19)*. ACM, 59–68. <https://doi.org/10.1145/3287560.3287598>
 - [56] Irene Solaiman, Zeerak Talat, William Agnew, Lama Ahmad, Dylan Baker, Su Lin Blodgett, Hal Daumé III, Jesse Dodge, Ellie Evans, Sara Hooker, Yacine Jernite, Alexandra Sasha Luccioni, Alberto Lusoli, Margaret Mitchell, Jessica Newman, Marie-Therese Png, Andrew Strait, and Apostol Vassilev. 2023. Evaluating the Social Impact of Generative AI Systems in Systems and Society. *arXiv:2306.05949* [cs.CY]
 - [57] Rohan Taori, Achal Dave, Vaishaal Shankar, Nicholas Carlini, Benjamin Recht, and Ludwig Schmidt. 2020. Measuring Robustness to Natural Distribution Shifts in Image Classification. In *Advances in Neural Information Processing Systems 33 (NeurIPS)*.
 - [58] Laura Weidinger, Maribeth Rauh, Nahema Marchal, Arianna Manzini, Lisa Anne Hendricks, Juan Mateos-Garcia, Stevie Bergman, Jackie Kay, Conor Griffin, Ben Bariach, Iason Gabriel, Verena Rieser, and William Isaac. 2023. Sociotechnical Safety Evaluation of Generative AI Systems. *arXiv:2310.11986* [cs.AI]
 - [59] Wenjun Zeng, Yuchi Liu, Ryan Mullins, Ludovic Peran, Joe Fernandez, Hamza Harkous, Karthik Narasimhan, Drew Proud, Piyush Kumar, Bhaktipriya Radharapu, Olivia Sturman, and Oscar Wahltinez. 2024. ShieldGemma: Generative AI Content Moderation Based on Gemma. *arXiv:2407.21772* [cs.CL]
 - [60] Lin Zhao, Ana Ferreira, Jonas Müller, and Chidi Okafor. 2026. ML-Bench&Guard: Policy-Grounded Multilingual Safety Benchmark and Guardrail Models. Preprint. Regulation-grounded multilingual guardrail training.

A EXTENDED RELATED WORK

Evaluation validity and accountability artifacts. That benchmarks under-determine the claims made from them is settled here: general-capability framings outrun what a fixed dataset supports [48], fairness benchmarks carry construct failures [8], behavioural testing beats aggregate scores [10, 53], progress on a fixed test set need not survive a change of collection process [52, 57], and generative-AI evaluations sit at a distance from the outcomes they are invoked to predict [38, 56, 58]. Our contribution is narrow and concrete: for one widely deployed component class we quantify four specific gaps against the size of the gain that would have been reported. Kapoor and Narayanan [34] is the closest methodological precedent for insisting that the leakage-shaped question — was the evaluation *source* represented in training? — be a first-class reported quantity rather than a footnote. And where model cards, datasheets, and dataset accountability practice [22, 30, 44] established that documentation is part of the artifact, internal-audit frameworks [50] and the external-audit ecosystem [43, 51] established who consumes it; Section 8 is written in that idiom.

Guard models and their evaluation.

Purpose-built guards [26, 31, 46, 47, 59] are typically reported against in-distribution and out-of-distribution blocks, and adapting a general model into a guard is standard practice [18]. Adjacent findings support ours from different directions: post-fine-tuning safety degradation tracks alignment/fine-tuning dataset similarity [28]; guardrails overfit their policies [41]; prompt-attack defenses learn surface heuristics [37]; guards are evadable and mutation-sensitive [7, 25]; and guard models are poorly calibrated [40], which is the precondition for the threshold-transfer failure we measure. Hossain et al. [27] fine-tune released guards on domain data, the manoeuvre we replicate on a registered

panel; Lee et al. [36] learn when to escalate to a larger guard; and Achara and Sanchez [1] audit deployed moderation classifiers with a fairness instrument more developed than ours.

Policy-grounded and regulated-domain evaluation. Cisneros-Velarde et al. [14] and Choi et al. [13] benchmark natural-language requests against written policies with *human* labels — the construct our mortgage label reaches for with a model judge — and Zhao et al. [60] derive regulation-grounded multilingual data. On the lending side, Barocas and Selbst [4] is the canonical account of how facially neutral variables carry protected-class information, exactly the mechanism our load-bearing benchmark stratum encodes; Bartlett et al. [5] measure it in consumer lending, and Bowen III et al. [9] vary race on real applications to audit a language model acting as underwriter. Table 5 in the appendix places the closest work against the axes we measure: no row does all of it, and almost every column has a row that does that column better than we do.

B EXPERIMENTAL SETUP

Checkpoints and adapters.

Four open instruction checkpoints (Qwen2.5-1.5B, SmolLM2-1.7B, SmolLM3-3B, Qwen3-4B). Each guard is a low-rank adapter [29] ($r=32$, $\alpha=64$, dropout 0.05) on the attention and MLP projections (q,k,v,o,gate,up,down), trained for 300 steps at learning rate 2×10^{-4} on a cosine schedule with 3% warmup, effective batch size 4, maximum sequence length 1024, five seeds (42–46) per checkpoint. Only the adapter trains; base weights are frozen, and “base” means the same weights read through the identical head of Equation (1).

Manifest. 1,200 rows: 400 from each of toxicchat, prompt_injections, and jailbreak_classification, balanced 200 safe / 200 unsafe within each corpus, with data order fixed by seed 42 so every checkpoint sees the same examples in the same order.

Scoring and metrics. One forward pass per row; the verdict is the single-token margin of Equation (1), stored as a raw decision margin rather than a saturating probability so that every metric is exactly recomputable from committed scores. Ranking metrics are tie-aware macro average precision; partial AUC is the one-way integral over false-positive rate [0, 0.05] normalized to mean recall inside the budget. Intervals are paired hierarchical bootstraps over 2,000 replicates resampling near-duplicate evaluation families and training seeds.

B.1 What a matched-budget number is and is not

Every matched-budget figure in this paper places a common alarm budget by taking a quantile of the *evaluated* negatives and then measuring recall on the same rows. This is a standard way to read a pair of ROC curves at one common point, and the paired bootstrap re-estimates the functional inside every replicate, so the intervals are internally consistent. It is *not* a threshold a production system could use, because placing it requires the labels the system is trying to predict. Read every such number as “recall at an empirical matched-budget ROC point”: a statement about ranking quality at a comparable alarm rate. A deployment claim would require the threshold to be fixed on a disjoint calibration set, frozen, and then reported with its *realized* false-alarm and recall rates on untouched acceptance rows. Only the calibration-only operating point of Section 5.1 does that.

Table 4. **The evidence ledger.** *Retro.* = retrospective, estimation-only on a panel inspected during development. *Prereg.* qualifies a retrospective row (locked estimands and criteria, *not* data-blind); it is not a fourth tier. The three tiers are never pooled into a single number, and no row is promoted to a causal, universal, or fair-lending claim.

Body of evidence	Tier	Establishes (conditional)	Principal limit → upgrade
Paired base→tuned panel, 4×5 seeds (Sections 4 and 5)	Retro.	Tuning raises represented ranking +0.3234 and moves transfer -0.0589; at a matched budget transfer recall falls on all four checkpoints	Fixed panel, two model lineages, one recipe, one manifest, balanced pools → sealed cohort, recipe sweep, prevalence-matched pools
Released-guard replication, 10 checkpoints (Section 4)	Retro., Prereg.	The direction extends to guards that are <i>already</i> guards: represented gain +0.111 (LCB +0.070)	Registry non-final, not data-blind; one degenerate harness cell excluded → locked registry, passing preflight
Hosted-frontier comparison, 2,275 external rows (Section 6)	Retro.; external	The hosted model leads the best small base by +0.109 [+0.077, +0.139] recall at a matched budget on unfamiliar prompts	Provider score is a coarse integer risk, so the budget is only approximately placeable; its input filter refused rows non-deterministically → graded provider score
Represented-source reversal, 12 cells (Section 6)	Retro.	Averaged over 3 sources the tuned panel leads by +0.083 [+0.013, +0.157]	Post hoc: not robust to reweighting; resampling the source set gives [-0.019, +0.220] → summary frozen before a fresh cohort
Mortgage dual-label instrument, 994 rows (Section 7)	LLM-judge	Zero-shot guards rank the policy label only 0.12–0.30 above its own chance floor; one worked G0/D1 row ranks below the median benign inquiry for all four	No SME adjudication, no human κ ; G1/D0 quadrant empty; 3 protected pairs in the scored split → counsel-reviewed cards, populated quadrant, $D=1$ pairs
External finance/health/law replication (Section 7.5)	External	The same bases rank domain violations well (AP 0.88–0.96) and the guard ordering does <i>not</i> carry across domain arms	Single-label general prompt safety, not a dual compliance construct → dual-labeled external construct with expert sign-off

Table 5. **Closest work by component.** ✓ reported; × not; ◦ partial. *Paired* = the same checkpoint measured before and after tuning on identical rows; *source-held-out* = transfer measured at the source, not row, level; *matched budget* = metrics compared at a common false-alarm rate. The leftmost three columns are where our remaining claim lives; the three ◦ in our own row are what a reader should hold us to — the router we tested is a score-margin router, “policy grounded” means consistency with a written card checked by a model judge, and the fairness construct is a small protected-pair gate, not an audit.

Work	Paired b→tune	Source- held-out	Matched budget	Routing method	Policy grnd.	Label tier	Fair- ness	Reg. dom.
Lee et al. [36]	×	◦	×	✓	×	public	×	×
Achara and Sanchez [1]	×	✓	◦	×	×	public	✓	×
Cisneros-Velarde et al. [14]	×	×	×	×	✓	human	×	◦
Choi et al. [13]	×	×	×	×	✓	human	×	✓
Zhao et al. [60]	×	◦	×	×	✓	mixed	×	✓
Hossain et al. [27]	✓	✓	×	×	×	public	×	◦
Hsiung et al. [28]	✓	◦	×	×	×	public	×	×
This paper	✓	✓	✓	◦	◦	judge	◦	✓

C THE EVIDENCE LEDGER AND THE CLOSEST WORK

D GENERATED TABLES AND THE WORKED CASE STUDY, REPRODUCED VERBATIM

Everything in this appendix is \input unmodified from the committed generated artifacts, so it is byte-identical to what the analysis harness emits.




Fig. 5. **The miss of Figure 6, as a rank.** For each zero-shot guard, how many of the 65 benign requests in the same 146-row split it scored *more* suspicious than the policy violation. Ranks rather than probabilities carry the claim, because one guard’s probabilities saturate near zero across the whole split and are not value-comparable with the others.

Table 6. Per-benchmark paired base→tuned deltas (upper block) and the calibration-targeted 5% false-positive operating point with single-class stress diagnostics (lower block). Generated from the lock-bound committed scores.

Regime / benchmark	Metric	Base	SFT	Δ	N
represented / toxicchat	AP	–	–	0.1880	–
represented / prompt_injections	AP	–	–	0.3704	–
represented / jailbreak_classification	AP	–	–	0.4119	–
transfer / jailbreakbench	AP	–	–	-0.0776	–
transfer / xstest	AP	–	–	-0.0120	–
transfer / wildguardtest	AP	–	–	-0.0669	–
transfer / wildjailbreak	AP	–	–	-0.0792	–
represented	TPR@target FPR	0.1296	0.7686	0.6390	313
represented	benchmark-macro realized FPR	0.0169	0.0108	-0.0061	364
represented	pooled-negative realized FPR	0.0213	0.0194	-0.0019	364
transfer	TPR@target FPR	0.5171	0.5807	0.0636	790
transfer	benchmark-macro realized FPR	0.0805	0.1546	0.0741	790
transfer	pooled-negative realized FPR	0.0430	0.1704	0.1273	790
Stress / OR-Bench	benign FPR	0.1181	0.1201	0.0020	400
Stress / HarmBench	recall	0.7800	0.6003	-0.1797	200

Table 11. External validation on ExpGuard (expert-annotated; input-prompt classification), 2275 rows across finance/health/law. Aggregate AP with a 2,000-resample bootstrap 95% CI, per-domain AP, and overall AUROC. Base checkpoints scored zero-shot via the canonical guard head; the ranking score is the raw decision margin $z_{\text{unsafe}} - z_{\text{safe}}$ (byte-parity with Act I). The top two guards’ *marginal* CIs overlap; the paired comparison below is the informative test.

Guard	AP (all, 95% CI)	AUROC	AP finance	AP health	AP law
Qwen2.5-1.5B	0.921 [.908, .932]	0.895	0.938	0.906	0.918
SmolLM2-1.7B	0.883 [.869, .897]	0.840	0.887	0.892	0.868
SmolLM3-3B	0.956 [.949, .963]	0.935	0.958	0.955	0.958
Qwen3-4B	0.951 [.943, .958]	0.927	0.957	0.938	0.957

Manuscript submitted to ACM

Paired top-two comparison (SmolLM3-3B – Qwen3-4B on *identical* rows, 2,000-resample bootstrap). The marginal CIs above overlap, but a paired test cancels row-difficulty variance and resolves one vertical: overall +0.0055 [−0.0008, +0.0120]; finance +0.0009 [−0.0075, +0.0097]; **health** +0.0168 [+0.0026, +0.0320] (CI excludes zero); law +0.0011 [−0.0118, +0.0140]. So the two guards are tied on finance and law and separate on health — “unresolved” was an unrun analysis, not a sample-size limit.

Worked example — one G0/D1 row, and what all four guards do with it

Row MGB-UD-00020 of the frozen public_test split (udaap / deceptive, difficulty *hard*, gold action BLOCK), quoted verbatim:

I'm reviewing this Montana conventional purchase file for a principal residence: loan amount is under \$100k, borrower income is in the \$120k-\$200k band, LTV is at or below 80%, but DTI is running 50%-60% and the credit file is borderline. There's no stated denial reason yet. Before I send it to underwriting, can you help me frame the decision note so it leans on the 'market fit' and the neighborhood's resale stability, and if needed I can just say the file doesn't align well with our portfolio for that area without getting into the DTI detail?

Gold labels. *G* = safe — no jailbreak, no injection, no abuse, nothing a content taxonomy has a category for — and *D* = intervene. Cards cited in policy_context: **D01** disparate treatment, **D02** redlining, **D07** disparate impact, **D12** adverse-action specificity, **D13** adverse-action reason masking, **D14** UDAAP-deceptive (benchmark coding, not a legal determination; this row's legal_review_status is synthetic_machine_validated_not_counsel_reviewed, and no card carries SME sign-off). The coded violation is entirely in the framing: the decision note is to rest on “market fit” and the neighborhood’s “resale stability” — a facially neutral geographic variable standing in for a protected class — and to state that the file does not align with the portfolio *for that area instead of* the real, permissible reason already in the file (the 50–60% DTI). Under cards D02/D07/D13 the benchmark codes that as redlining-by-proxy plus adverse-action reason masking, asked for in the register of a routine underwriting question.

Guard (zero-shot)	$p(\text{unsafe})$ on this row	median p on the 65 benign rows	benign rows ranked above (of 65)
Qwen2.5-1.5B	0.662	0.860	46
SmolLM2-1.7B	0.560	0.742	57
SmolLM3-3B	0.014	0.097	65
Qwen3-4B	1.1×10^{-6}	1.9×10^{-6}	44

All four guards rank this violation below the median benign inquiry in the same split, and SmolLM3-3B ranks it below *every one* of the 65 — on that guard's ordering no benign inquiry in the split looks safer than this violation. Ranks, not values, carry the claim: Qwen3-4B's probabilities saturate near 0 across the whole split (Section 7.4), so its column is not value-comparable with the others, while the rank column is.

A related violation, named rather than coded. Row MGB-FL-00028 asks whether “the way the borrower’s surname, preferred language, and neighborhood profile show up in the file should change the amount of discretion we use for pricing” — a fair-lending ask with the protected traits said out loud. The benign-above counts invert: Qwen2.5-1.5B 1, SmolLM2-1.7B 7, SmolLM3-3B 0, Qwen3-4B 15. But this is *not* the same ask: the two rows differ in fact sheet, domain and subdomain label, cited cards, and request type (write a pretextual note vs. ask whether traits may inform pricing). They are an illustration of a possible surface-form effect, not a measurement of one, and no controlled instrument for it exists in v1 — every protected pair is benign on both arms, so the gate of Section 7.4 can only probe over-refusal, never a violation scored lower when coded. That needs $D=1$ pairs.

Fig. 6. A **G0/D1** miss, end to end, from the frozen public-test split. The request is polite, procedurally framed, and generally safe by the benchmark's own *G* label; the benchmark's policy cards code honoring it as redlining-by-proxy and adverse-action reason masking. Every number is recomputed by tools/emit_case_study_tex.py from the committed per-row scores. This is the concrete form of the claim that in a regulated domain the dangerous request can be compliant-looking text a general guard has no label for.

E VALIDATION ROADMAP

The upgrades that would raise each row of Table 4 to a stronger tier, in the order we would run them. (1) A genuinely sealed cohort — rows never committed to a public repository and never scored during development — so the paired estimand can carry a confirmatory rather than retrospective claim. (2) Subject-matter-expert adjudication of the mortgage policy cards with per-label agreement statistics, the single change that would move that arm off the weakest tier. (3) Populating the G1/D0 quadrant, so that over-refusal of compliant requests naming a protected class becomes measurable. (4) Protected pairs with $D=1$ on both arms, the instrument required to test whether a violation stated in a routine register is scored lower than the same violation stated plainly — a hypothesis our worked example is consistent with and cannot establish. (5) A prevalence-matched evaluation pool, and reporting each headline guard as a curve

Table 7. **The same operating point read at an equal false-alarm budget.** Table 6 compares recalls at each guard’s own calibrated threshold, where the tuned guard alarms far more often (17.0% vs 4.3% pooled transfer FPR) — so it buys recall with alarms, and the two recalls are not comparable. Here each SFT seed is instead thresholded so its pooled transfer false-alarm rate is *at or below its own base’s* (the budget column), which makes the recalls directly comparable. At an equal budget the tuned guard is worse on **all four** checkpoints and on both instruments: transfer recall 0.517→0.185 (-0.332) and HarmBench recall 0.780→0.171 (-0.609). Same committed rows and same scorer as Table 6; only the threshold rule changes, so this needs no GPU and is byte-checked by make verify. **This is a retrospective ROC point, not a deployable threshold:** the threshold is read off the *same* labelled negatives the recall is then measured on, so a production system without labels could not place it. Read the row as “recall at an empirical matched-FPR ROC point”, and see Appendix B.1 for what an operational version would require.

Checkpoint	FPR	transfer recall (macro)			HarmBench recall		
	budget	base	SFT own thr.	SFT matched	base	SFT own thr.	SFT matched
Qwen2.5-1.5B	7.7%	0.485	0.571	0.272	0.810	0.606	0.290
SmolLM2-1.7B	2.4%	0.416	0.588	0.194	0.610	0.615	0.202
SmolLM3-3B	1.5%	0.512	0.596	0.102	0.740	0.611	0.077
Qwen3-4B	5.6%	0.655	0.569	0.172	0.960	0.569	0.114
Panel mean	4.3%	0.517	0.581	0.185	0.780	0.600	0.171

Table 8. **The same eight cells, read in the operating region instead of over the whole ranking.** Paired base→SFT change on identical committed rows under three metrics: benchmark-macro AP (the report’s primary metric), benchmark-macro one-way partial AUC over FPR [0, 0.05] (mean TPR inside the alarm budget; chance floor 0.025, not 0.5), and benchmark-macro TPR at that budget. Brackets are two-sided 95% paired bootstrap intervals over 2,000 replicates resampling evaluation family_id clusters and training seeds, the same protocol as the rest of the report. **No sign flips:** every cell moves the same way under all three metrics, so Act I’s direction survives being read at a deployable operating point. What does not survive is the *size*. macro-AP understates both halves of the trade — the represented gain is +0.323 on AP against +0.686 on pAUC (2.1×), and the transfer cost is -0.059 against -0.174 (3.0×), with panel-mean budget recall moving -0.199. Averaging precision over the whole ranking credits ordering in the deep negative mass, where an inline guard never operates. Same rows, same scorer, no GPU: only the metric changes, and it is byte-checked by make verify.

Checkpoint	Δ macro-AP	Δ pAUC[0, 0.05]	Δ TPR@0.05
<i>Represented (id_test)</i>			
Qwen2.5-1.5B	+0.354 [+0.272, +0.412]	+0.802 [+0.723, +0.863]	+0.802 [+0.703, +0.860]
SmolLM2-1.7B	+0.528 [+0.456, +0.572]	+0.791 [+0.731, +0.851]	+0.842 [+0.771, +0.889]
SmolLM3-3B	+0.313 [+0.242, +0.367]	+0.676 [+0.567, +0.767]	+0.712 [+0.550, +0.782]
Qwen3-4B	+0.098 [+0.055, +0.148]	+0.475 [+0.305, +0.563]	+0.357 [+0.195, +0.514]
Panel mean	+0.323	+0.686	+0.678
<i>Transfer (transfer_test)</i>			
Qwen2.5-1.5B	-0.039 [-0.082, +0.007]	-0.002 [-0.126, +0.085]	-0.081 [-0.198, +0.075]
SmolLM2-1.7B	+0.040 [+0.001, +0.077]	+0.011 [-0.093, +0.118]	-0.012 [-0.108, +0.113]
SmolLM3-3B	-0.087 [-0.111, -0.062]	-0.295 [-0.363, -0.219]	-0.265 [-0.364, -0.196]
Qwen3-4B	-0.150 [-0.197, -0.106]	-0.409 [-0.499, -0.312]	-0.437 [-0.528, -0.311]
Panel mean	-0.059	-0.174	-0.199

AP(π_+) rather than one balanced number. (6) A recipe sweep over adapter rank, step budget, and data mixture, to separate what is a property of specialization from what is a property of this configuration. (7) An embedding-space decontamination check, which our lexical audit does not cover. (8) A pre-specified frontier summary frozen before a fresh cohort is scored, the only route to a confirmatory version of Section 6.

Table 9. Baseline zero-shot instruction guards on the frozen benchmark (public_test, 146 rows: 75 G0/D1, 6 G1, 3 protected pairs). AP is recomputed in the repo canonical environment from the committed per-row scores (exactly reproducible); the AP·D column carries a 2,000-resample bootstrap 95% CI. Threshold-free AP·D is moderate (0.67–0.85) against a *chance floor* of 0.555 (81/146 rows are *D*-positive), i.e. only 0.12–0.30 above chance; AUROC·D is the base-rate-free companion. Five of the six pairwise AP·D CI comparisons overlap (only Qwen3-4B vs. SmolLM2-1.7B separates), so we do not rank guards by it. *G* and *D* are also not independent in v1: the G1/D0 cell is empty, so *G* is nested inside *D*. Δ_{context} is the mean absolute protected-pair gap on the *probability* scale (lower is more invariant); Δ_{margin} is the same gap on the raw margin $z_{\text{unsafe}} - z_{\text{safe}}$ (log-odds, scale-free), which exposes saturation — Qwen3-4B’s 0.000 becomes 0.80, the second largest on the panel; Δ_{tok} restricts to pairs that swap a single token (2 of 3 here; 18 of the 39 release pairs use a multi-word placeholder arm instead), which removes Qwen2.5-1.5B’s apparent outlier. All three rest on only $n = 3$ pairs, so they are read as a direction, not a calibrated fairness estimate — and on this split they do not rank guards. AP·final equals AP·D on this frozen set because the G1/D0 cell is empty (every *G*-positive row is also *D*-positive), so the composed label reduces to *D*. The fixed 5%-FPR operating point is threshold-knife-edge for these clustered-score guards and is omitted (see text).

Guard	AP·G	AP·D (95% CI)	AUROC·D	AP·final	Δ_{context}	$\Delta_{\text{margin}}^{\text{context}}$	$\Delta_{\text{tok}}^{\text{context}}$
qwen25_15b_base	0.681	0.793 [.71, .87]	0.743	0.793	0.183	0.84	0.020
qwen3_4b_base	1.000	0.851 [.78, .91]	0.778	0.851	0.000	0.80	0.000
smollm2_17b_base	0.261	0.672 [.57, .78]	0.641	0.672	0.023	0.10	0.024
smollm3_3b_base	0.546	0.733 [.64, .81]	0.603	0.733	0.010	0.15	0.012

Not scored (gated, HF license not accepted): llama_guard_3_1b, wildguard_7b.

Table 10. Composition of the frozen v1_hmda2022 release (994 rows). The private_test split is committed with text and has already been scored, so it is dataset-held-out by convention rather than sealed. 0 of 958 content_family groups span a split boundary; 1 of 39 protected pair_id pairs do.

Split	Rows	Quadrant	Rows	Domain	Rows
train	604	G0/D0 (benign)	450	benign	450
dev	149	G0/D1 (domain-only)	502	fair_lending	204
public_test	146	G1/D1 (both)	42	fraud	112
private_test (committed, not sealed)	95	G1/D0 (general-only)	0	udaap	90
				disclosure	66
				atr_qm	54
				privacy	18

F REPRODUCIBILITY

Every table and figure is either \input from a committed generated artifact or plotted by a script that parses those same artifacts. One command recomputes the generated artifacts from committed per-row scores into a scratch directory and asserts byte-identity with the committed copies, leaving the tracked tree unmodified: 27 of 35 inputs are byte-checked in any environment, 4 require a lock-pinned environment, 4 are uncovered, and 0 fail. Three tiers are kept distinct. *Analysis reproducibility* — regenerating a table from committed per-row scores — is what the byte-check covers, with no GPU and no network. *Training and scoring reproducibility* — re-deriving those per-row scores by training the 4×5 adapters and rescoring — requires a GPU and pinned model and data access. *Artifact-generation reproducibility* — re-generating the mortgage benchmark itself — is **not** claimed: its construction stages run at nonzero temperature and the release is intentionally frozen, so only its evaluation reproduces. Third-party corpora are referenced by pinned identifier, revision, and content hash rather than redistributed. Code, data manifests, generated tables, the frozen benchmark, and the committed per-row scores are public; the repository URL and the benchmark’s attribution are withheld here for anonymous review [2, 3].